

Retrieval-Augmented Generation

RAG combines information retrieval with text generation. A system retrieves relevant document passages before asking a language model to generate an answer.

Document Processing

RAG systems commonly load documents, split them into chunks, create embeddings, and store the embeddings in a vector database.

Embeddings

An embedding represents text as a numerical vector. Similar meanings can be located near each other in vector space.

Vector Database

A vector database stores embeddings and supports similarity search to find relevant chunks.

Advantages

RAG can help a language model answer questions using a specific document collection and can reduce the need to place every document directly in the prompt.